

Parallel Architectures and Algorithms for Large-Scale Nonlinear Programming

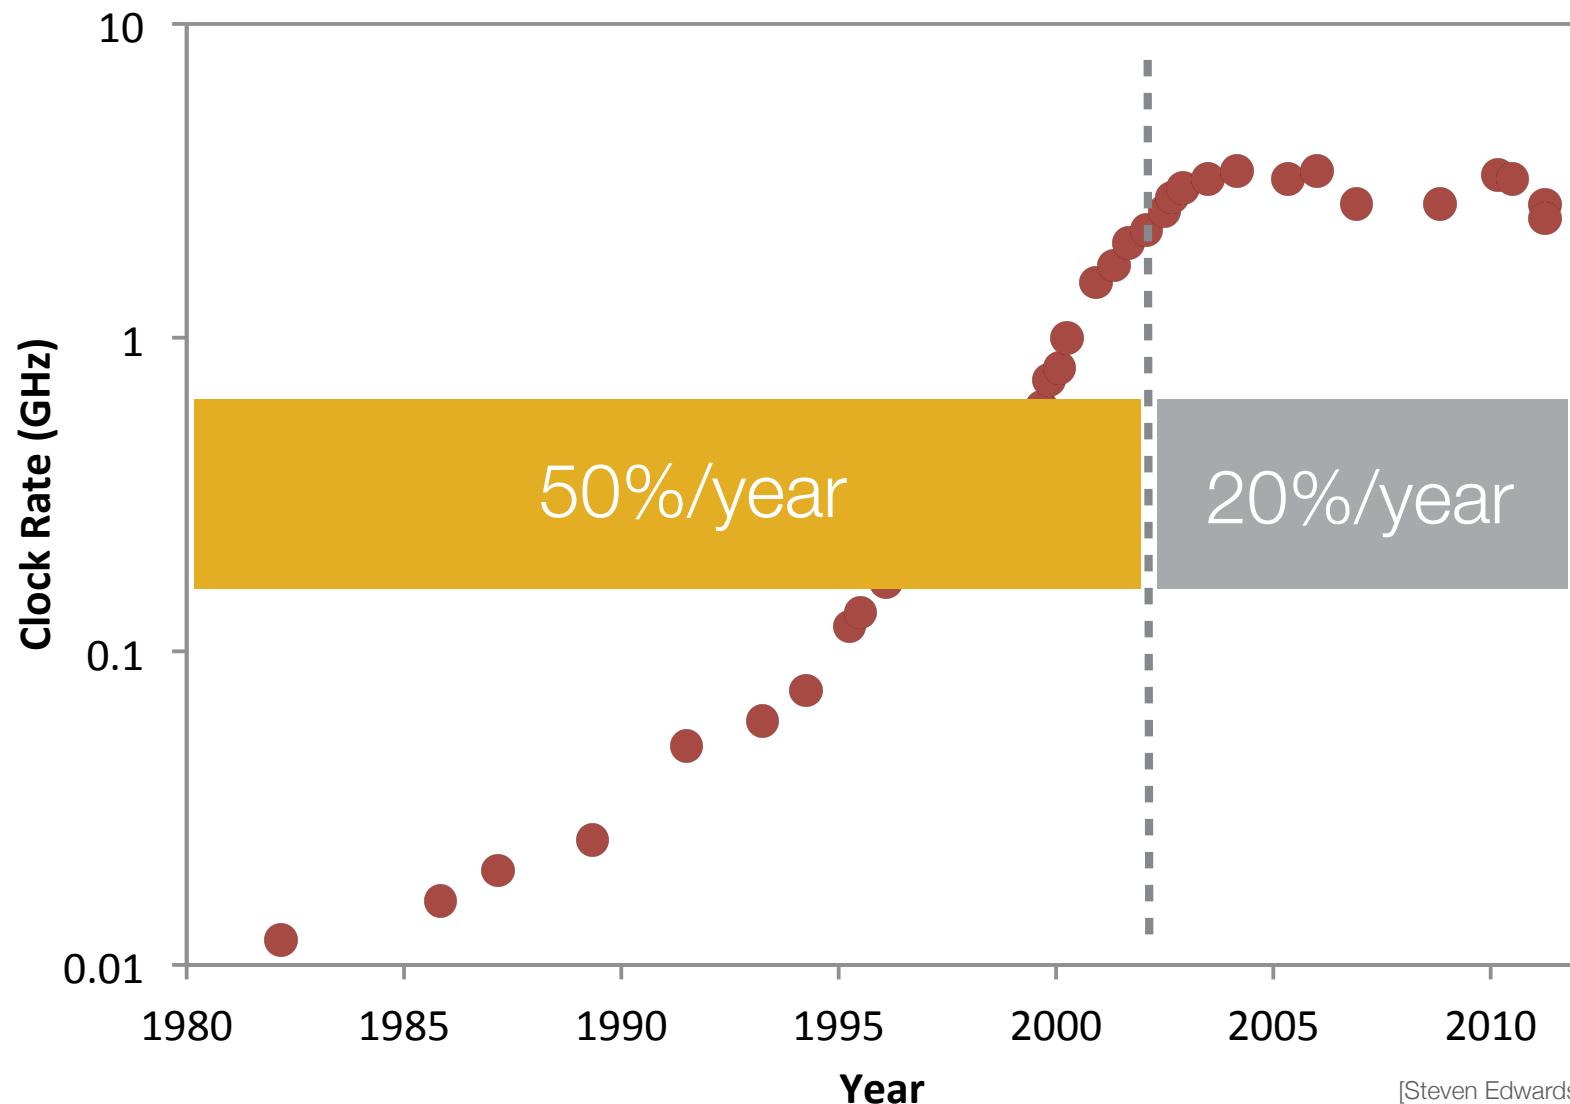
Carl D. Laird

Associate Professor, School of Chemical Engineering, Purdue University
Faculty Fellow, Mary Kay O'Connor Process Safety Center



Laird Research Group: <http://allthingsoptimal.com>

Landscape of Scientific Computing

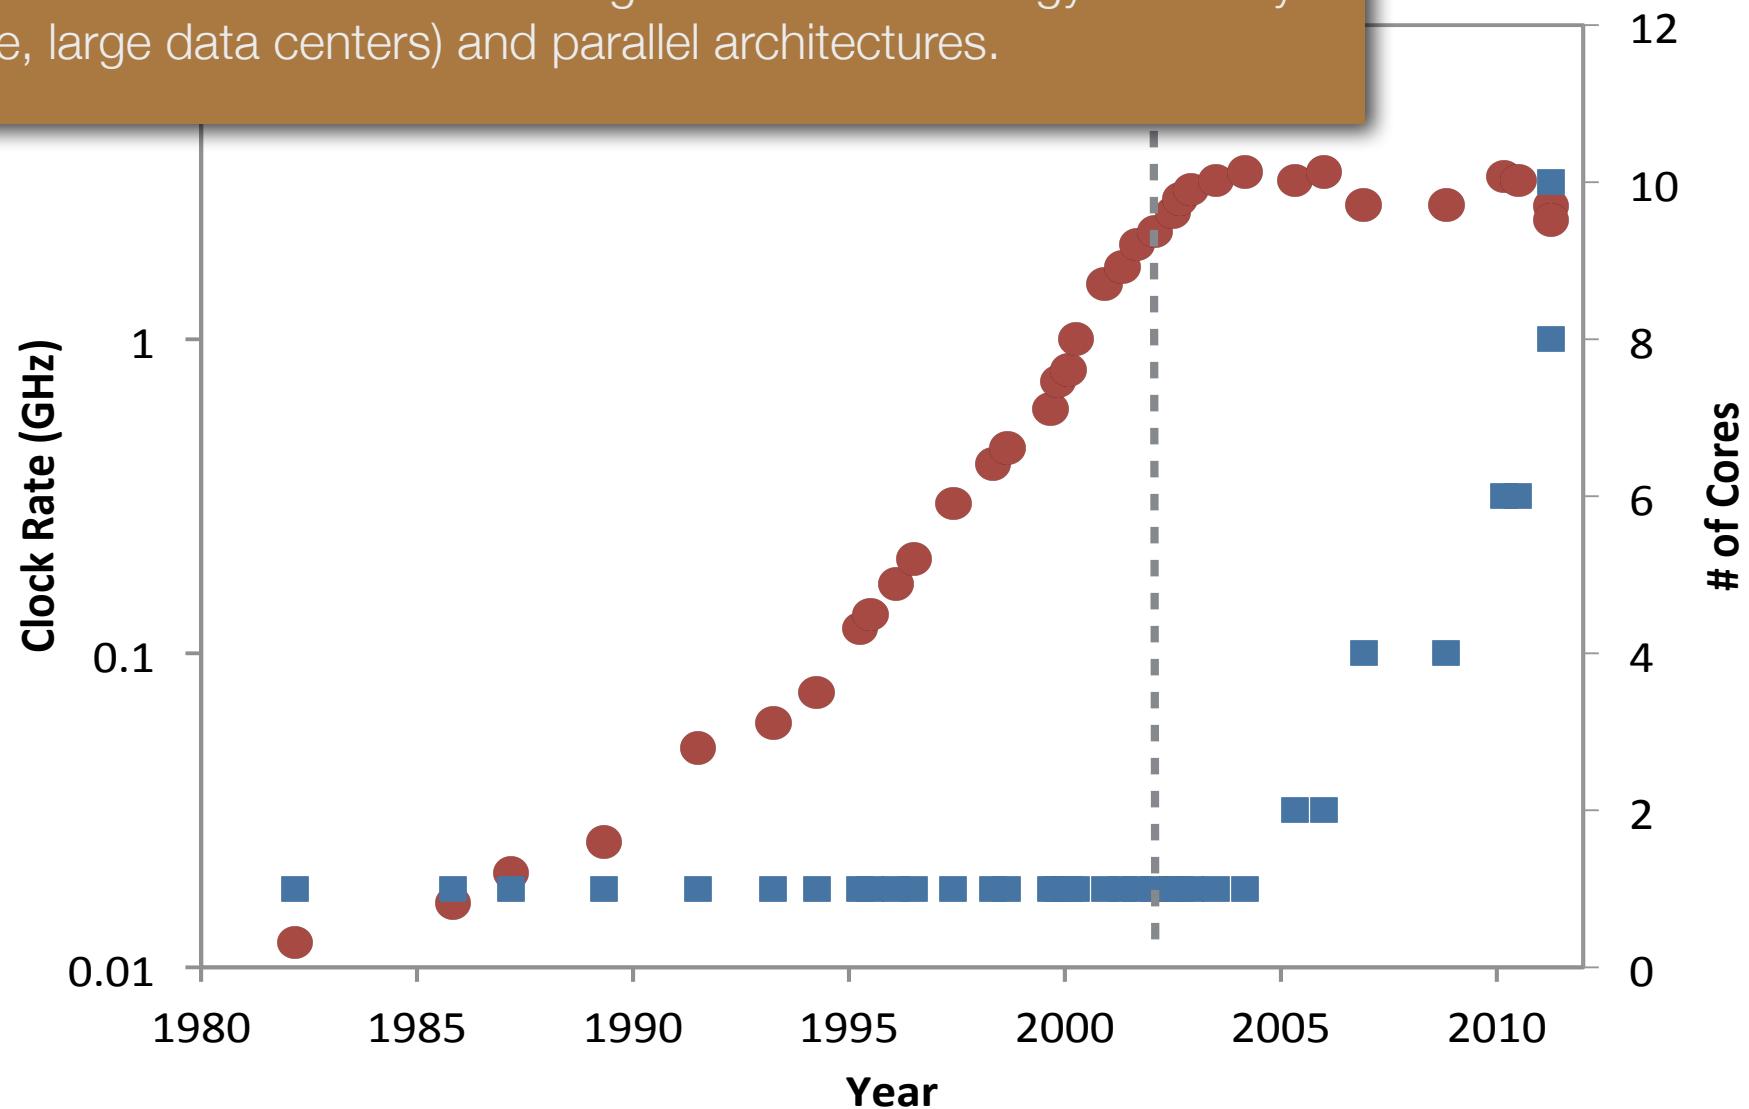


[Steven Edwards, Columbia University]

Landscape of Scientific Computing

Clock-rate, the source of past speed improvements have stagnated.

Hardware manufacturers are shifting their focus to energy efficiency (mobile, large data centers) and parallel architectures.



Frontiers in Chemical Engineering: Research Needs and Opportunities (1988)

<http://www.nap.edu/openbook/030903793X/html/137.html>, copyright 1988, 2000 The National Academy of Sciences, all rights reserved

Panel on Computer Assisted Process and Control Engineering

ARTHUR W. WESTERBERG (*Chairman*),
Carnegie Mellon University

HENRY CHIEN, Monsanto Corporation

JAMES M. DOUGLAS, University of Massachusetts

BRUCE A. FINLAYSON, University of Washington

ROLAND KEUNINGS, University of California, Berkeley

MANFRED MORARI, California Institute of Technology

JEFFREY J. SIROLA, Eastman Kodak Company

WILLIAM SILLIMAN, Exxon Production Research Company

USING THE COMPUTER'S POTENTIAL

Each decade over the last 35 years has seen the processing speed of newly designed computers increase by a factor of about 100 owing to advances in the design of electronic microcircuits and other computer hardware. On top of this has been another 100-fold increase per decade in computer speed, thanks to more efficient methods of carrying out computations (algorithms). It is not widely appreciated that new algorithms have been as valuable as hardware design in improving computer performance. With the combination of improved computer hardware and better algorithms, effective computer speeds have more than doubled on average each year.

“... over a 15-year span... [problem] calculations improved by a factor of 43 million. [A] factor of roughly 1,000 was attributable to faster processor speeds, ... [yet] a factor of 43,000 was due to improvements in the efficiency of software algorithms.”
- attributed to Martin Grotschel

[Steve Lohr, “Software Progress Beats Moore’s Law”, New York Times, March 7, 2011]

Landscape of Computing



Landscape of Computing



High-Performance
Parallel Architectures

Multi-core, Grid,
HPC clusters,
Specialized
Architectures (GPU)

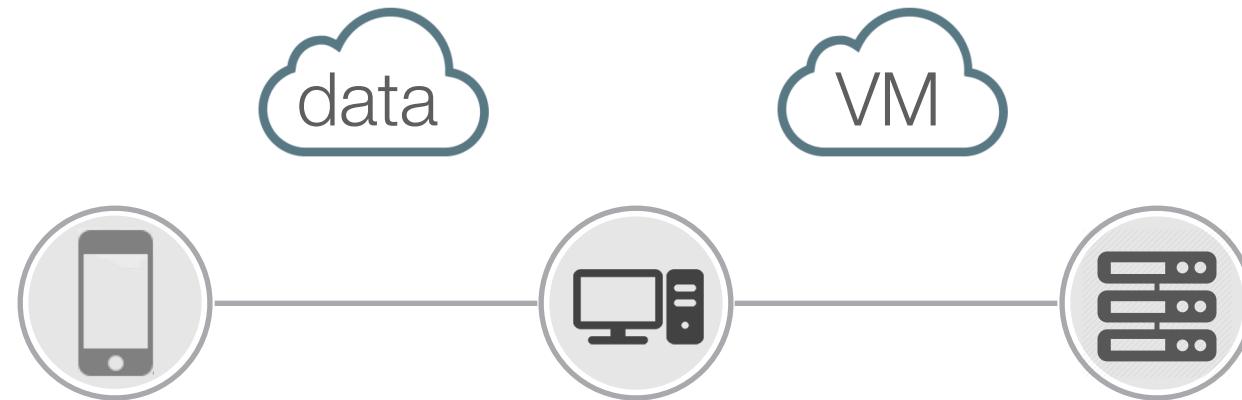
Landscape of Computing



High-Performance
Parallel Architectures

Multi-core, Grid,
HPC clusters,
Specialized
Architectures (GPU)

Landscape of Computing



iOS and Android device sales have surpassed PC

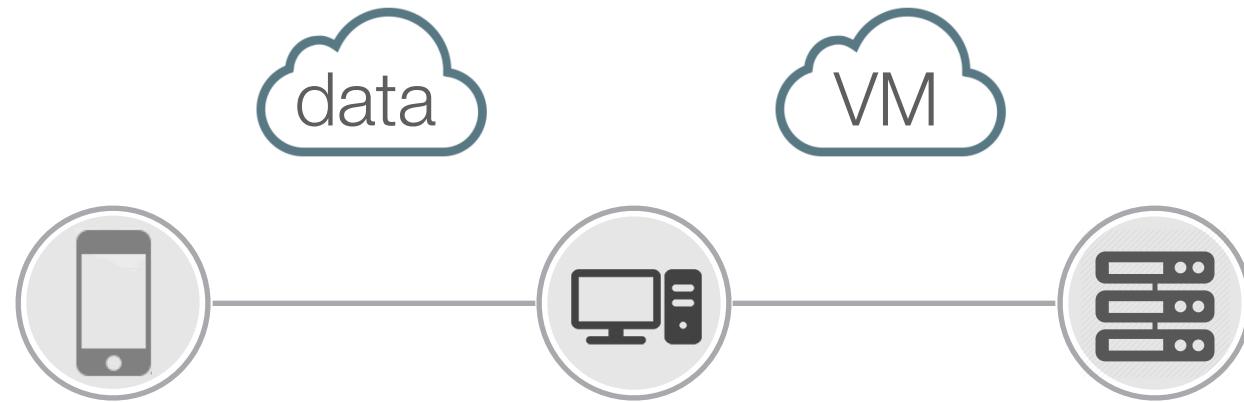
iPhone performance
~2X / year

iPhone 5S within a factor of 2 of 2.5 GHz Core i5 [Geekbench 3]

High-Performance Parallel Architectures

Multi-core, Grid, HPC clusters, Specialized Architectures (GPU)

Landscape of Computing



Watch: increased use in operations

High-Performance Parallel Architectures

Watch: rise of medical devices and health

Multi-core, Grid, HPC clusters, Specialized Architectures (GPU)

Watch: touch form factor for desktop computing

High-Performance Computing



Parallel Computing Architectures

Measuring Parallel Performance

Parallel Algorithms for NLP

Parallel Computing Architectures

Flynn's Taxonomy

	Single Data	Multiple Data
Single Instruction	SISD	SIMD
Multiple Instruction	MISD	MIMD

Parallel Computing Architectures

Flynn's Taxonomy

	Single Data	Multiple Data
Single Instruction	SISD	SIMD
Multiple Instruction	MISD	MIMD

Beowulf Cluster (MIMD)

- Scalable Distributed Computing
100s of Processors
- Standard Compilers (MPI)
- Network Communication (Ethernet)
- Communication Bottleneck

Desktop Multi-core (MIMD)

- Affordable Hardware
Low Number of Processors
- Standard Compilers (Threads, OpenMP)
- Fast Communication (No Network)
- Memory Access/# CPU Bottleneck

Parallel Computing Architectures

Flynn's Taxonomy

	Single Data	Multiple Data
Single Instruction	SISD	SIMD
Multiple Instruction	MISD	MIMD

Emerging Architectures (SIMD/Hybrid)

Graphics Processing Unit (GPU)

- Affordable, 1000's of Processors
- Specialized Compilers/Languages
 - CUDA, Brook++, OpenCL
- Several Complexities & Limitations

Beowulf Cluster (MIMD)

- Scalable Distributed Computing
100s of Processors
- Standard Compilers (MPI)
- Network Communication (Ethernet)
- Communication Bottleneck

Desktop Multi-core (MIMD)

- Affordable Hardware
Low Number of Processors
- Standard Compilers (Threads, OpenMP)
- Fast Communication (No Network)
- Memory Access/# CPU Bottleneck

Writing High-Performance Scientific Software

- Implementation and hardware details matter
 - Cache sizes have a major impact
 - Memory latency, coalescing, and alignment matter
 - Network latency and bandwidth matter
 - Compiler optimization can only do so much
 - Use high-performance code tailored to architecture (e.g. BLAS & LAPACK for dense linear algebra)
- Languages matter
 - C / C++ / Fortran still the dominant choices
 - Interpreted languages (E.g. Matlab / Python) might be reasonable, but be careful
 - Modern compiled languages close to speed of C (with significant reduction in development time)

Measuring Parallel Performance

Speedup: $S = \frac{\text{Time for best known serial algorithm}}{\text{Time for parallel algorithm}}$

Efficiency: $\frac{\text{Actual Speedup}}{\text{Ideal Speedup}} = \frac{S}{N_p}$ usually < 1

Communication Overhead

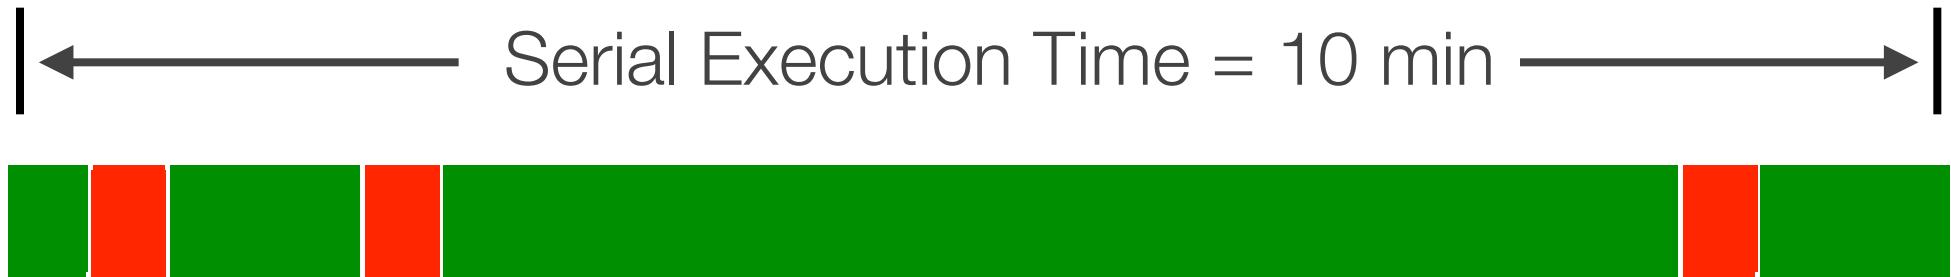
Memory/Cache Bottlenecks

Measuring Parallel Performance

| ← → | Serial Execution Time = 10 min |



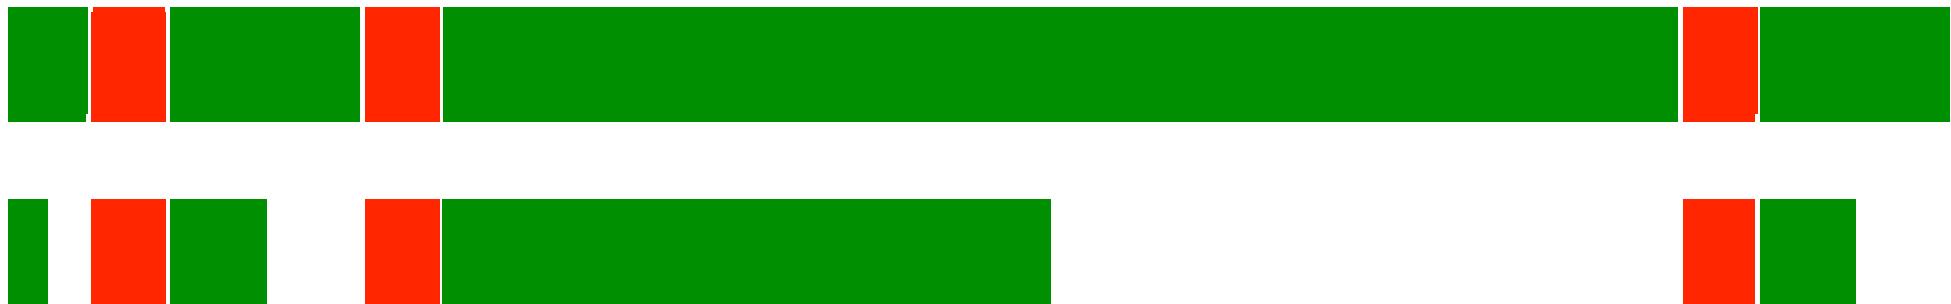
Measuring Parallel Performance



- Can be computed in parallel
- Must be computed in serial

Measuring Parallel Performance

| ← → | Serial Execution Time = 10 min | ← → |



■ Can be computed in parallel

■ Must be computed in serial

Measuring Parallel Performance

| ← → | Serial Execution Time = 10 min | → |



| ← → | Parallel Time = 5.5 min | → |

2 Processors
1.8 Times Speedup
~90% Efficiency

■ Can be computed in parallel

■ Must be computed in serial

Measuring Parallel Performance

| ← → | Serial Execution Time = 10 min | → |



| ← → | Parallel Time = 5.5 min | → |
2 Processors
1.8 Times Speedup
~90% Efficiency

A horizontal bar divided into 10 equal segments. The first two segments are red, and the remaining eight are green. A double-headed arrow above the bar indicates it is being processed by 2 processors simultaneously, resulting in a parallel execution time of 5.5 minutes.

Infinite Processors
10 Times Speedup

■ Can be computed in parallel ■ Must be computed in serial

Measuring Parallel Performance

| ← → | Serial Execution Time = 10 min | → |



Amdahl's Law, Maximum Speedup: $S_{\infty} = \frac{1}{\phi_s}$



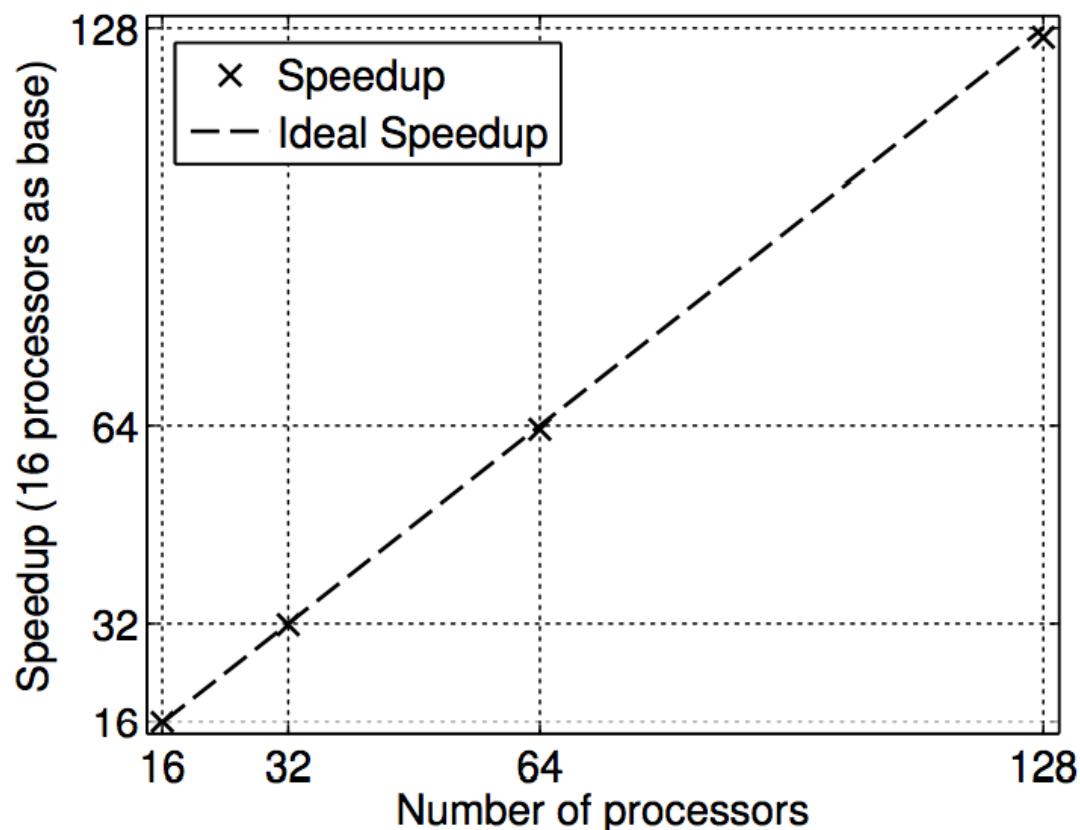
| ← → | Parallel Time = 1.0 min |

Infinite Processors
10 Times Speedup

■ Can be computed in parallel

■ Must be computed in serial

Measuring Parallel Performance (Scaling Efficiency)



Strong Scaling

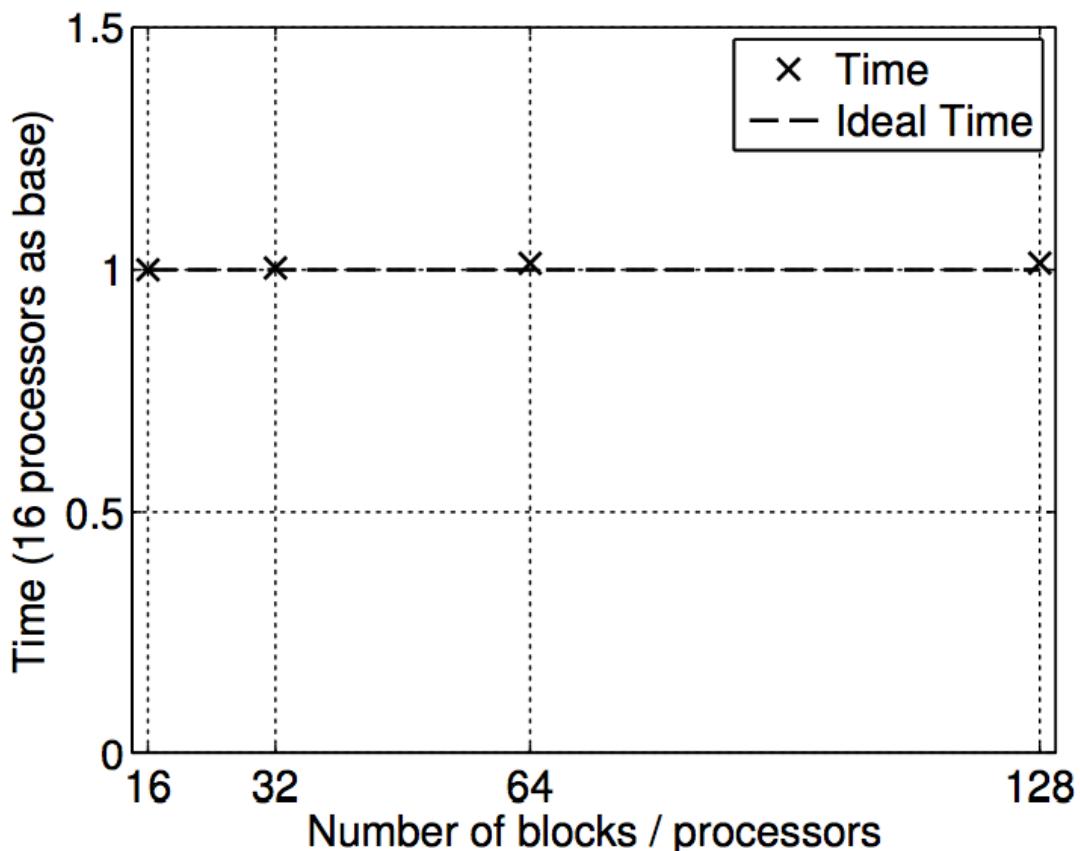
- Problem size fixed as # processors increased
- Fixed workload

Linear Scaling

- Speedup = # Processors

Performance compared against linear scaling

Measuring Parallel Performance (Scaling Efficiency)



Weak Scaling

- Problem size grows as # processors increases
- Fixed workload per processor

Linear Scaling

- Constant run-time

Performance compared against linear scaling

Parallel Algorithms for Nonlinear Optimization



Algorithm Parallel on
General Problem

Often based on less
effective serial algorithms

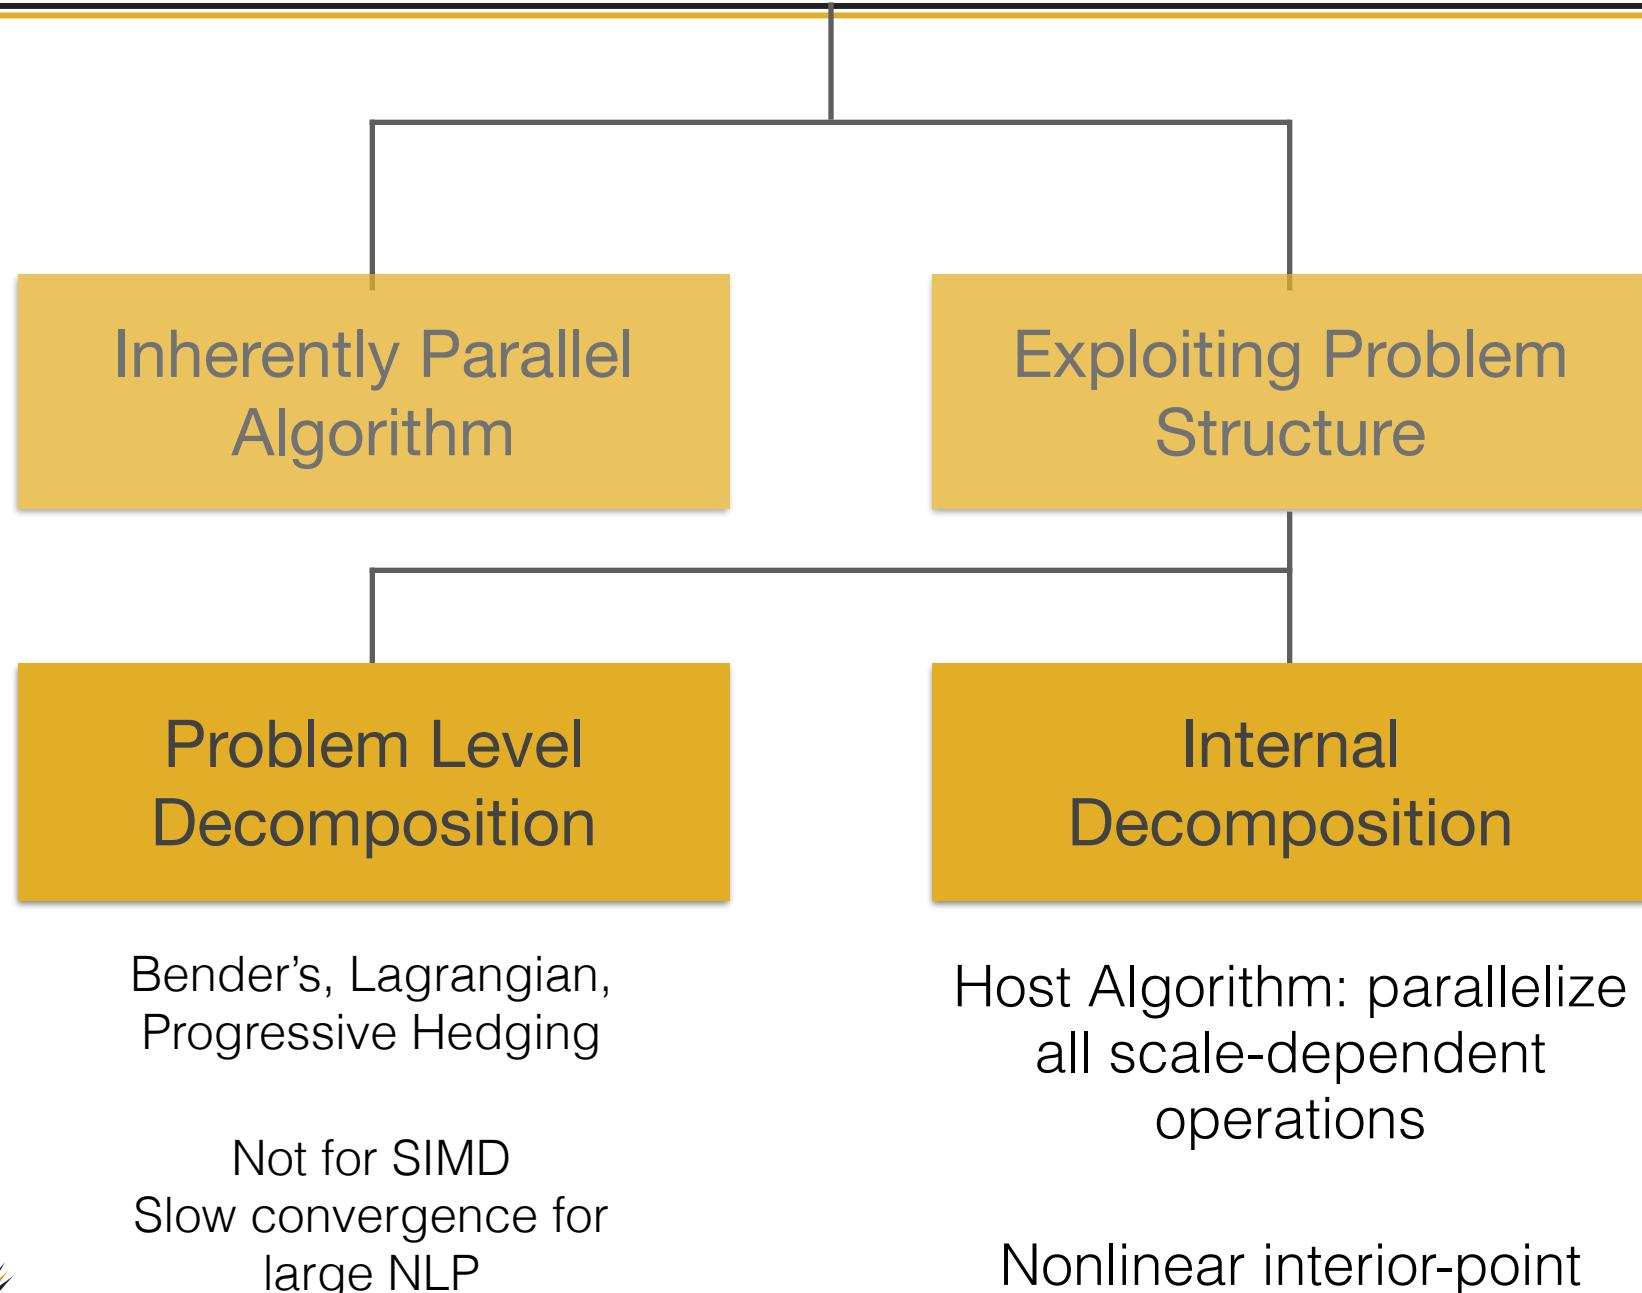
Easily parallelized

Algorithm Parallel Due to
Problem Structure

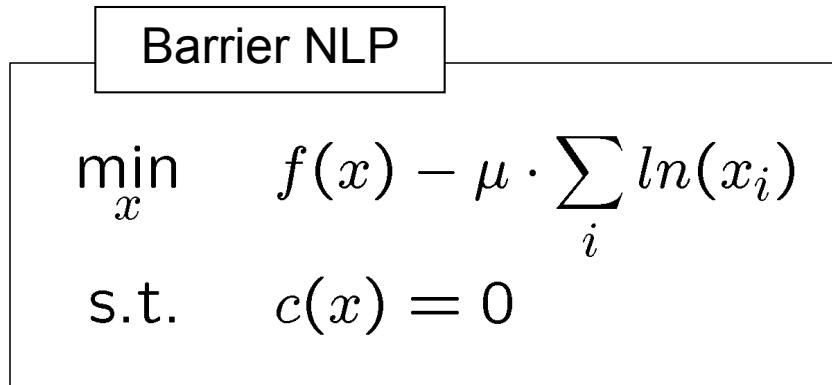
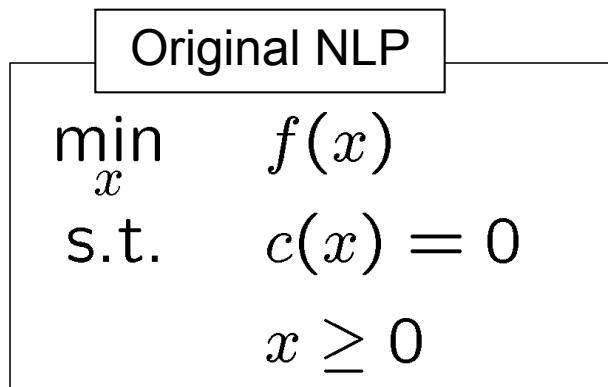
Decomposition
strategies

High-level parallelism

Parallel Algorithms for Nonlinear Optimization



Nonlinear Interior-point Methods



as $x \rightarrow 0, -\ln(x) \rightarrow \infty$

as $\mu \rightarrow 0, x^*(\mu) \rightarrow x^*$

Fiacco & McCormick (1968)

Nonlinear Interior-point Methods

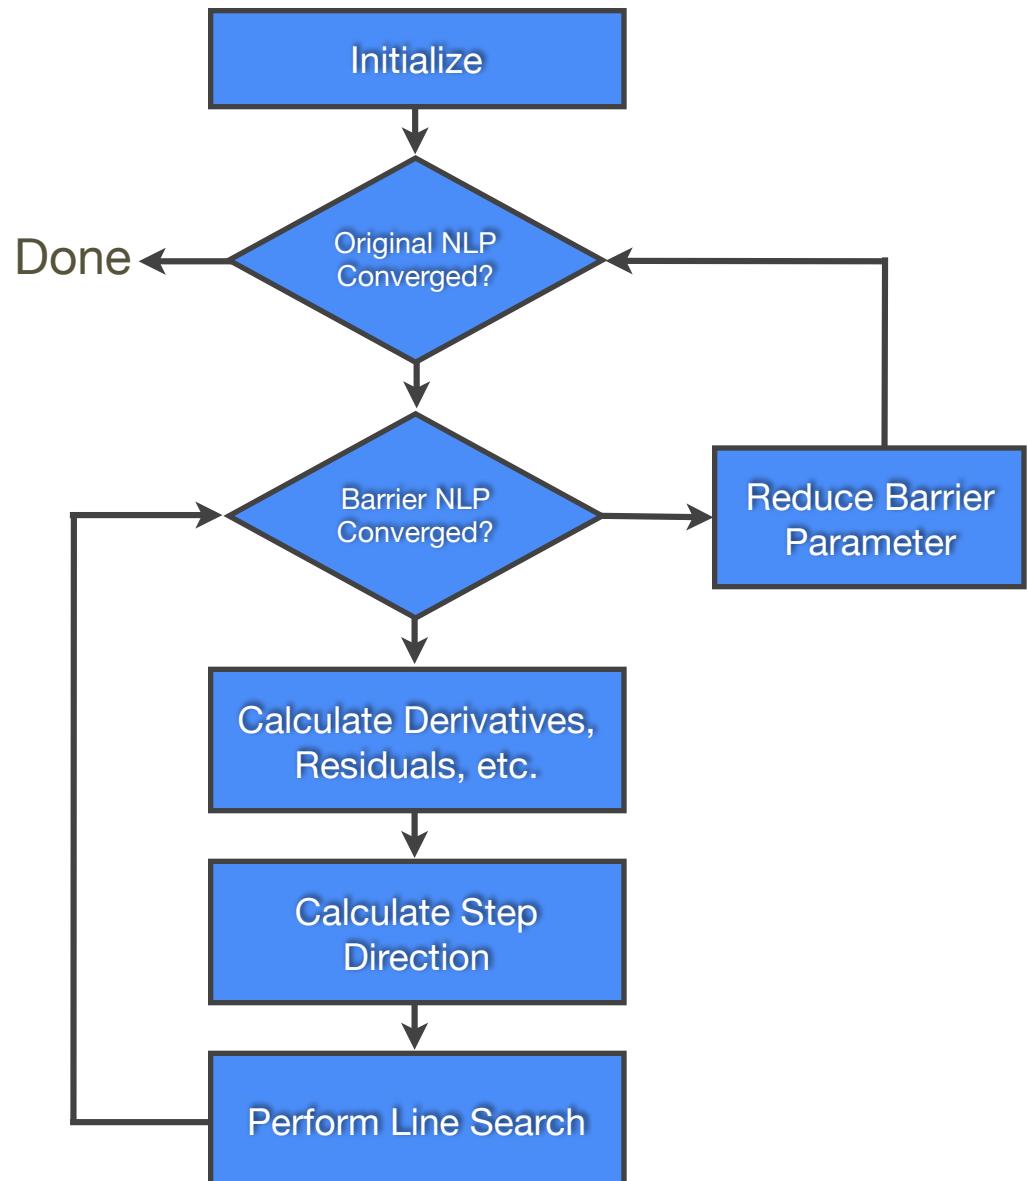
Original NLP

$$\begin{aligned} \min_x \quad & f(x) \\ \text{s.t.} \quad & c(x) = 0 \\ & x \geq 0 \end{aligned}$$

Barrier NLP

$$\begin{aligned} \min_x \quad & f(x) - \mu \cdot \sum_i \ln(x_i) \\ \text{s.t.} \quad & c(x) = 0 \end{aligned}$$

KNITRO (Byrd, Nocedal, Hribar, Waltz)
LOQO (Benson, Vanderbei, Shanno)
IPOPT (Wachter, Laird, Biegler)



Nonlinear Interior-point Methods

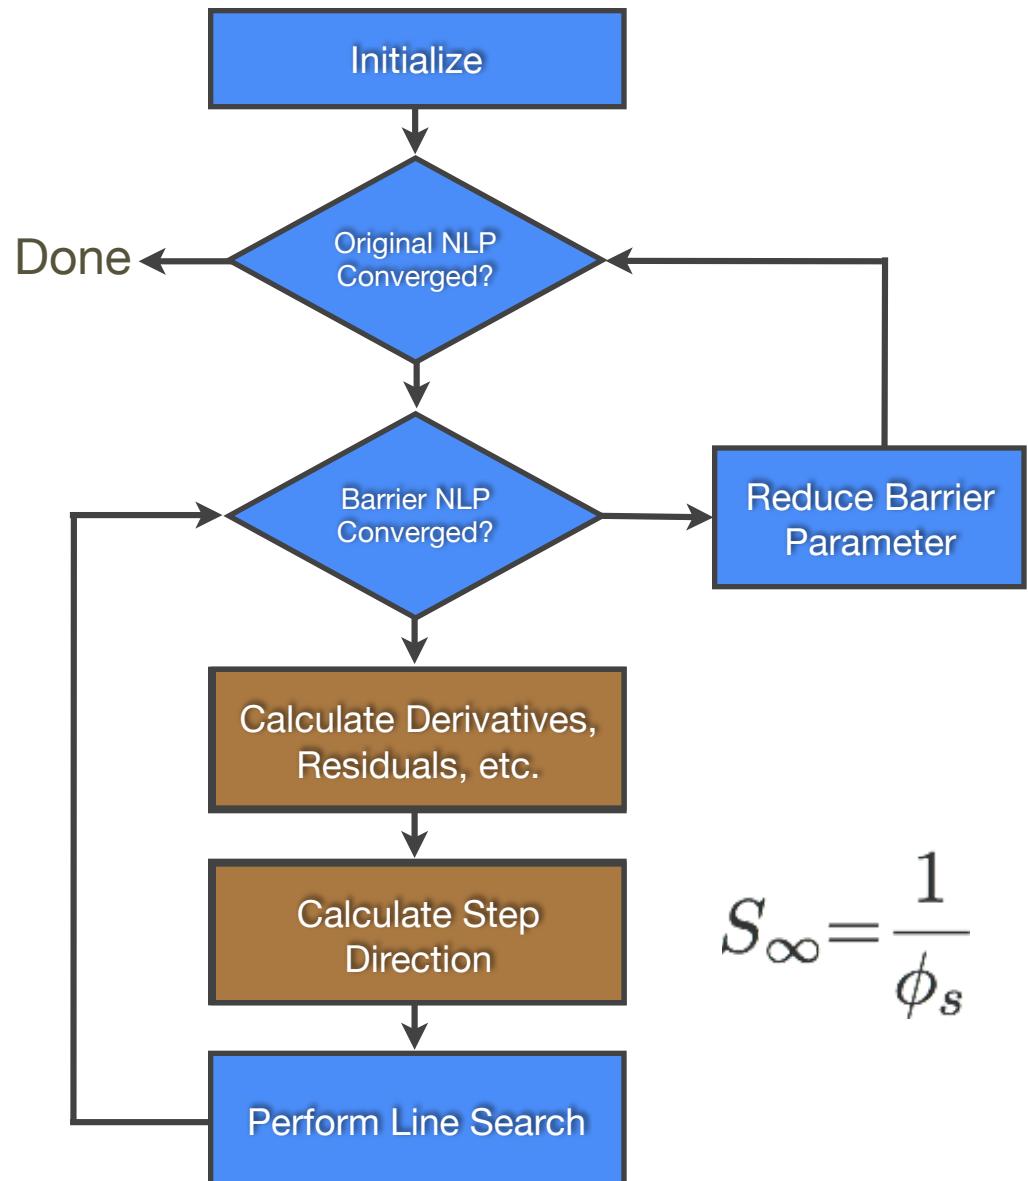
Original NLP

$$\begin{aligned} \min_x \quad & f(x) \\ \text{s.t.} \quad & c(x) = 0 \\ & x \geq 0 \end{aligned}$$

Barrier NLP

$$\begin{aligned} \min_x \quad & f(x) - \mu \cdot \sum_i \ln(x_i) \\ \text{s.t.} \quad & c(x) = 0 \end{aligned}$$

KNITRO (Byrd, Nocedal, Hribar, Waltz)
LOQO (Benson, Vanderbei, Shanno)
IPOPT (Wachter, Laird, Biegler)



$$S_{\infty} = \frac{1}{\phi_s}$$

Parallel Nonlinear Interior-point Methods

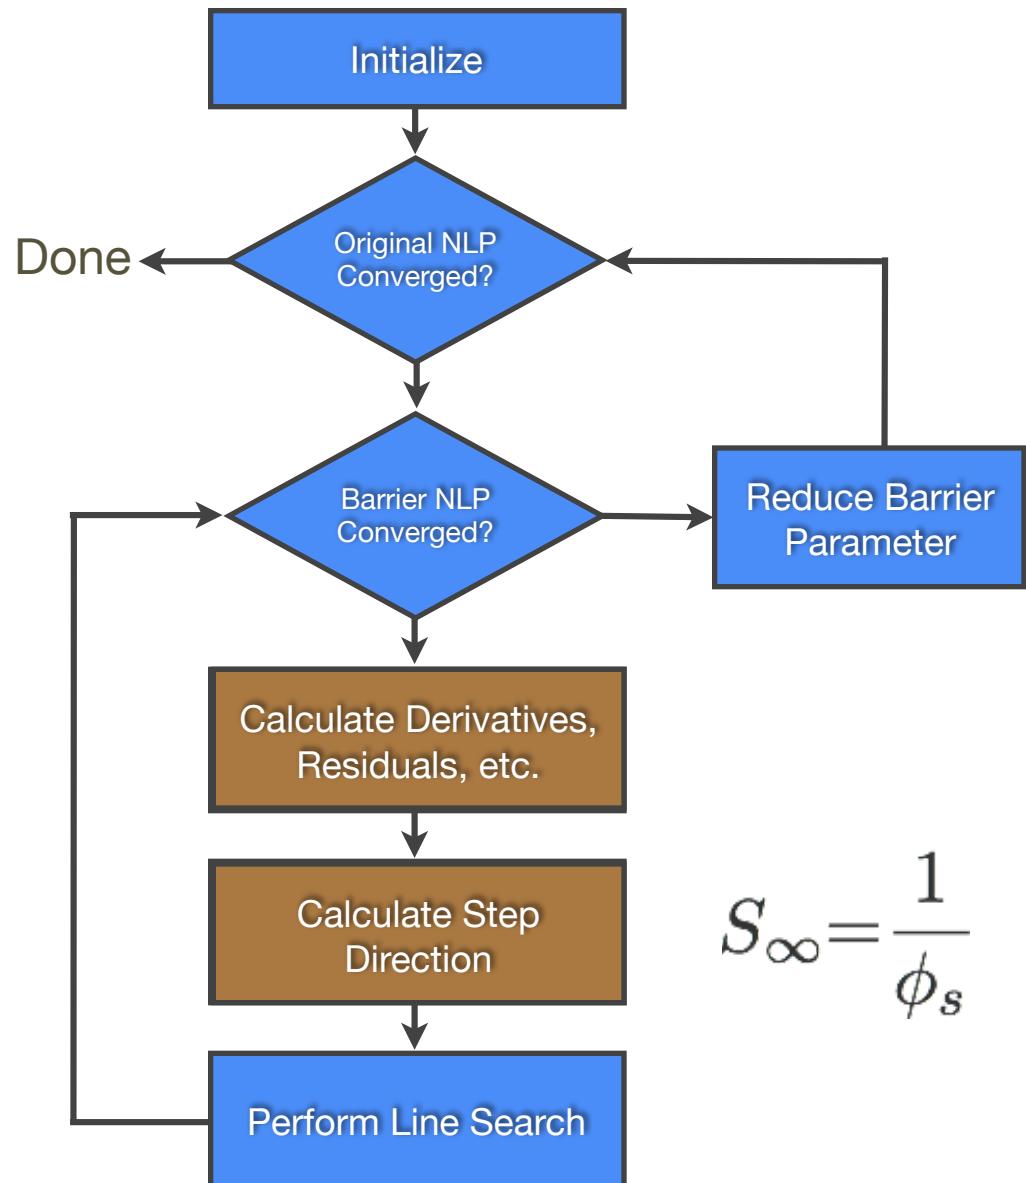
Structure in the optimization problem induces structure in the linear algebra

Parallelize all scale-dependent operations

- Vector and matrix operations
- Model evaluation

Compared with problem-level decomposition, implementation is time consuming

Retain convergence properties of serial algorithm



$$S_{\infty} = \frac{1}{\phi_s}$$

Parallel Linear Algebra for NLP

General Purpose Parallel Direct Solvers

- For general sparse systems - no need to specify structure
- [Schenk & Gartner 2004; Scott 2003; Amestoy et al. 2000; ...]
- Modest scale up

Iterative Linear Solvers (e.g. GMRES, PCG)

- Capable of handling general sparse systems
- Easily parallelized (many implementations)
- Appropriate for SIMD architectures
- Requires effective preconditioning
[Dollar et al. 2007; Forsgren et al. 2008]

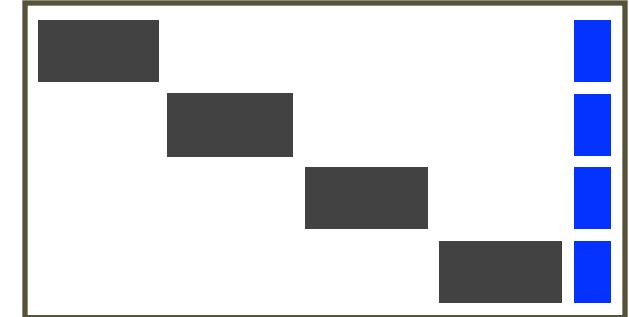
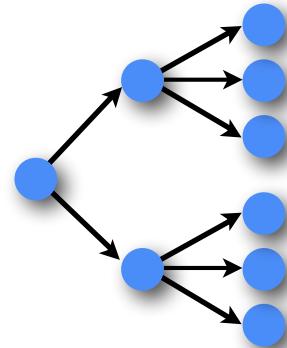
Tailored Decomposition of Problem Structure

- Custom parallel linear algebra exploits a fixed, known structure
- [Gondzio, Anitescu (PIPS), Laird, ...]
- Problem class specific - difficult to implement
- Can scale to large parallel architectures
(100's - 1000's of processors)

Exploiting Problem Structure

Optimization Under Uncertainty

- block structure because of coupled scenarios
- common structure of many applications

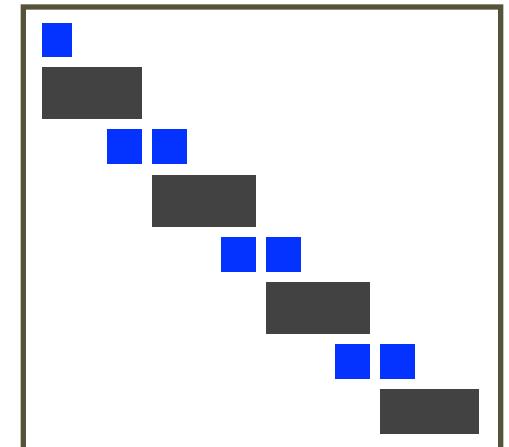


Kang, J., Word, D.P., and Laird, C.D., "An interior-point method for efficient solution of block-structured NLP problems using an implicit Schur-complement decomposition", to appear in Computers and Chemical Engineering, 2014.

Dynamic Optimization

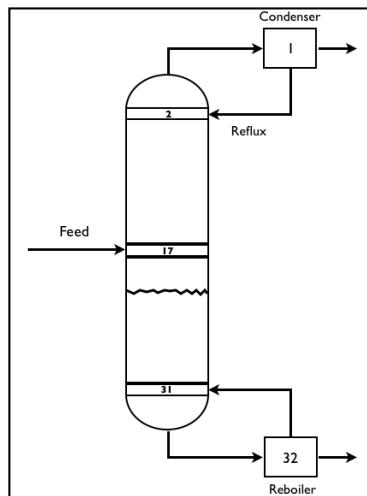
- block structure because of finite element discretization

$$\begin{aligned} & \min_u \int_{t_0}^{t_f} L(x, y, u) \, dt \\ \text{s.t. } & F(\dot{x}, x, y, u) = 0 \\ & x(t_0) = x_0 \\ & (x, y, u)^L \leq (x, y, u) \leq (x, y, u)^U \end{aligned}$$



Word, D.P., Kang, J., Akesson, J., and Laird, C.D., "Efficient Parallel Solution of Large-Scale Nonlinear Dynamic Optimization Problems", to appear in Computational Optimization and Applications, 2014.

Parallel Performance: Optimization Under Uncertainty



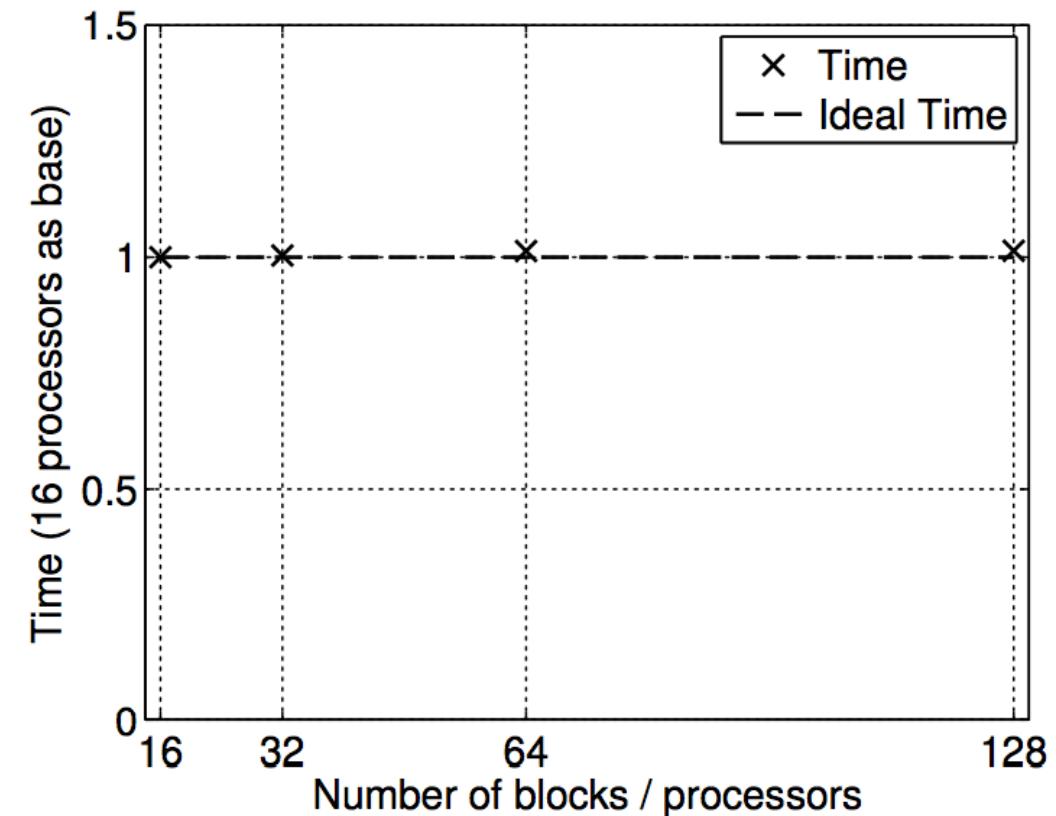
- 32 state variables, 35 algebraic variables
- Discretize model (OCFE)
- Uncertainty in mole fraction of the feed stream
- 96 scenarios, 32 processors

[Benallou, Seborg, and Mellichamp (1986)]

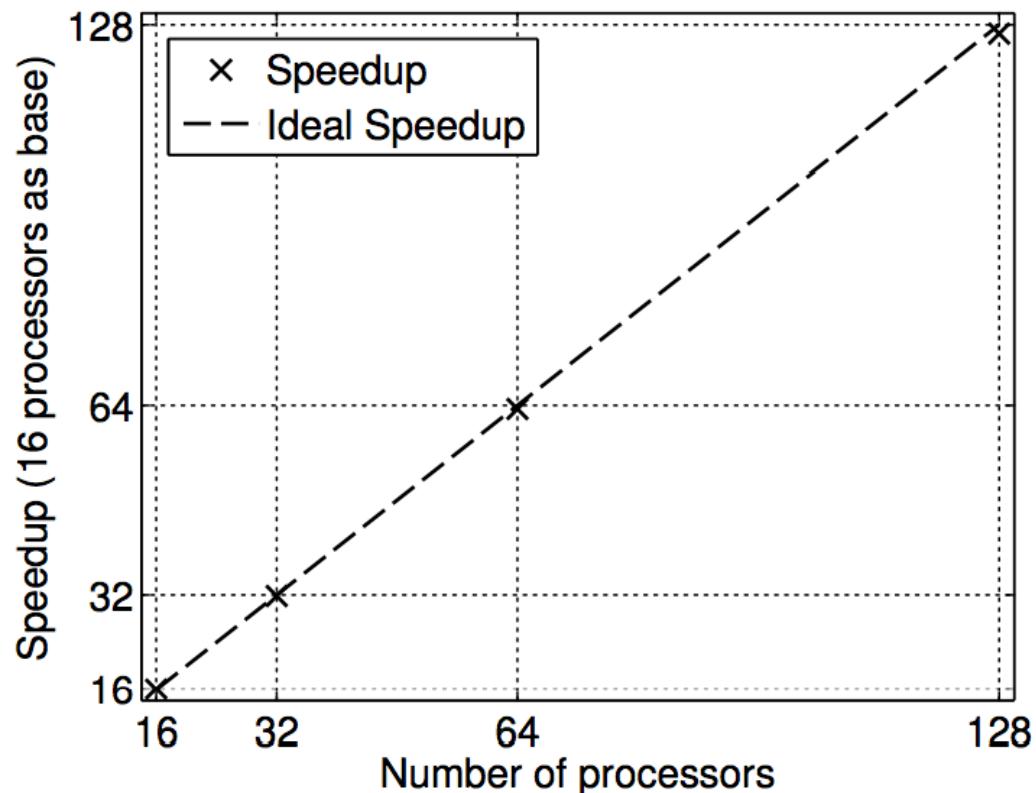
Case	# Vars.	# Coupling Vars.	FS-S time(s)	ESC-S time(s)	ESC-P time(s)	PCGSC-S time(s)	PCGSC-P time(s)
1	1430550	150	10.3	79.1	2.6	17.9	0.6
2	2861100	300	-	-	10.8	-	1.1
3	4291650	450	-	-	32.1	-	2.4
4	5722200	600	-	-	70.3	-	3.2
5	7152750	750	-	-	90.5	-	4.3
6	8583300	900	-	-	160.5	-	5.3
7	10013850	1050	-	-	218.0	-	6.3
8	11444400	1200	-	-	286.6	-	8.1

Kang, J., Word, D.P., and Laird, C.D., "An interior-point method for efficient solution of block-structured NLP problems using an implicit Schur-complement decomposition", to appear in Computers and Chemical Engineering, 2014.

Weak and Strong Scaling: Optimization Under Uncertainty



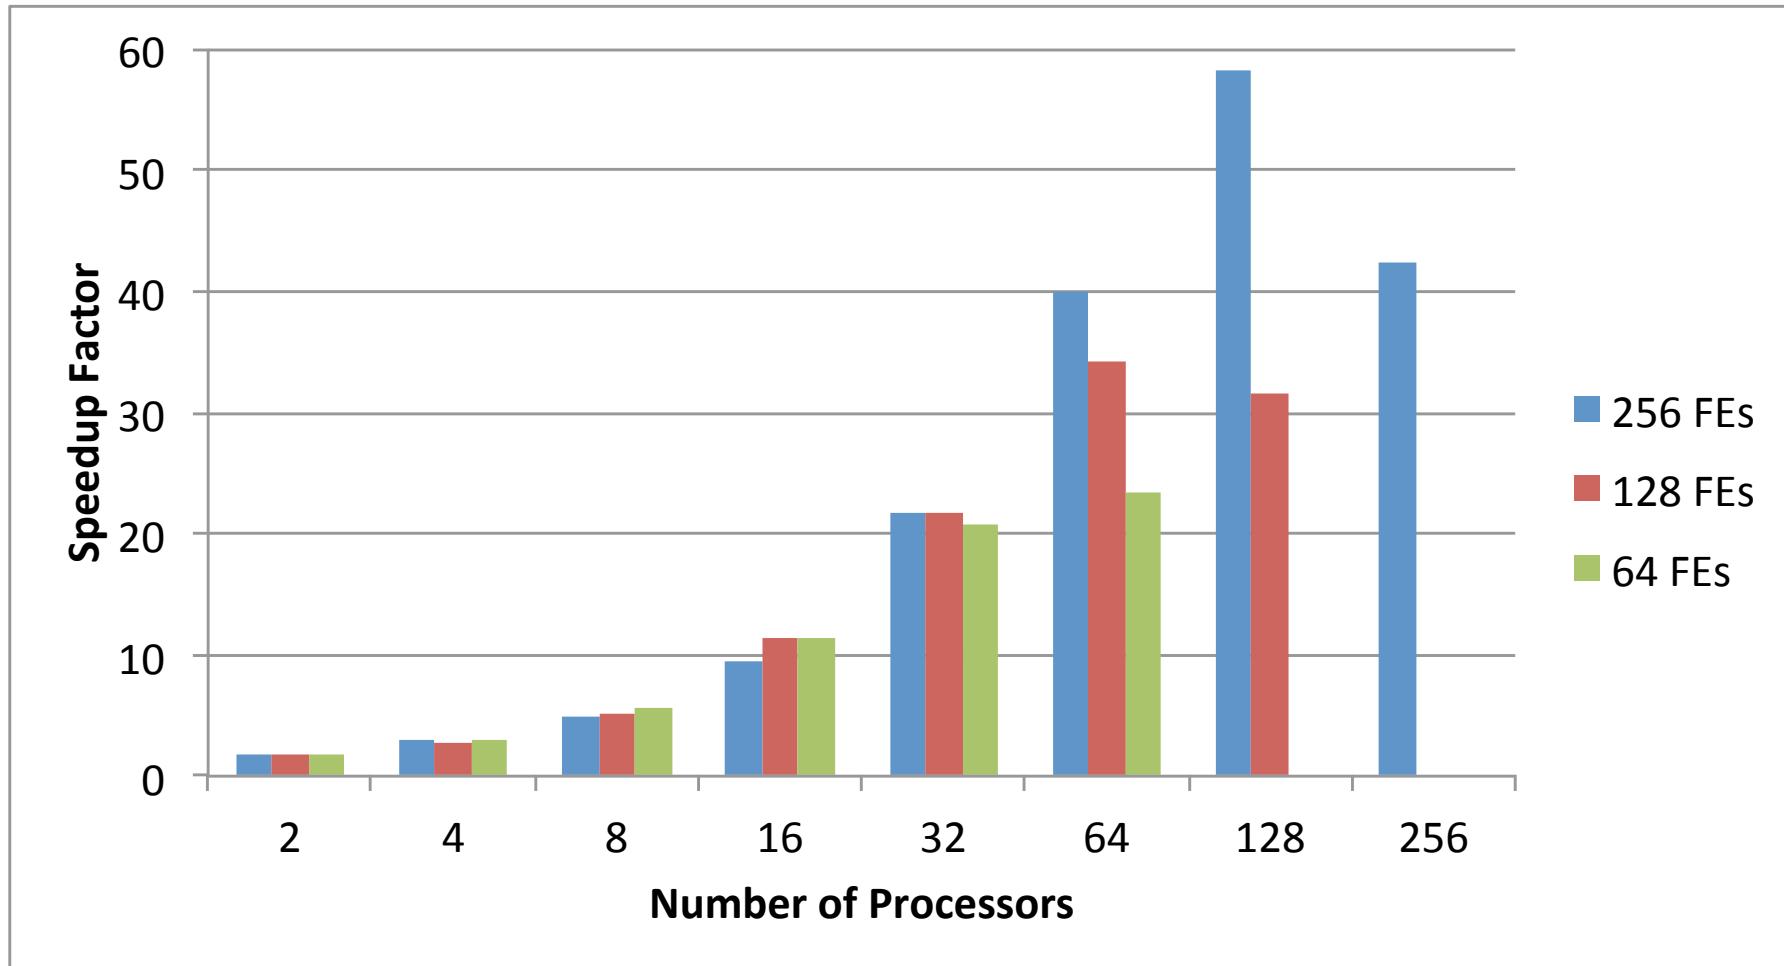
(a) Weak Scaling



(b) Strong Scaling

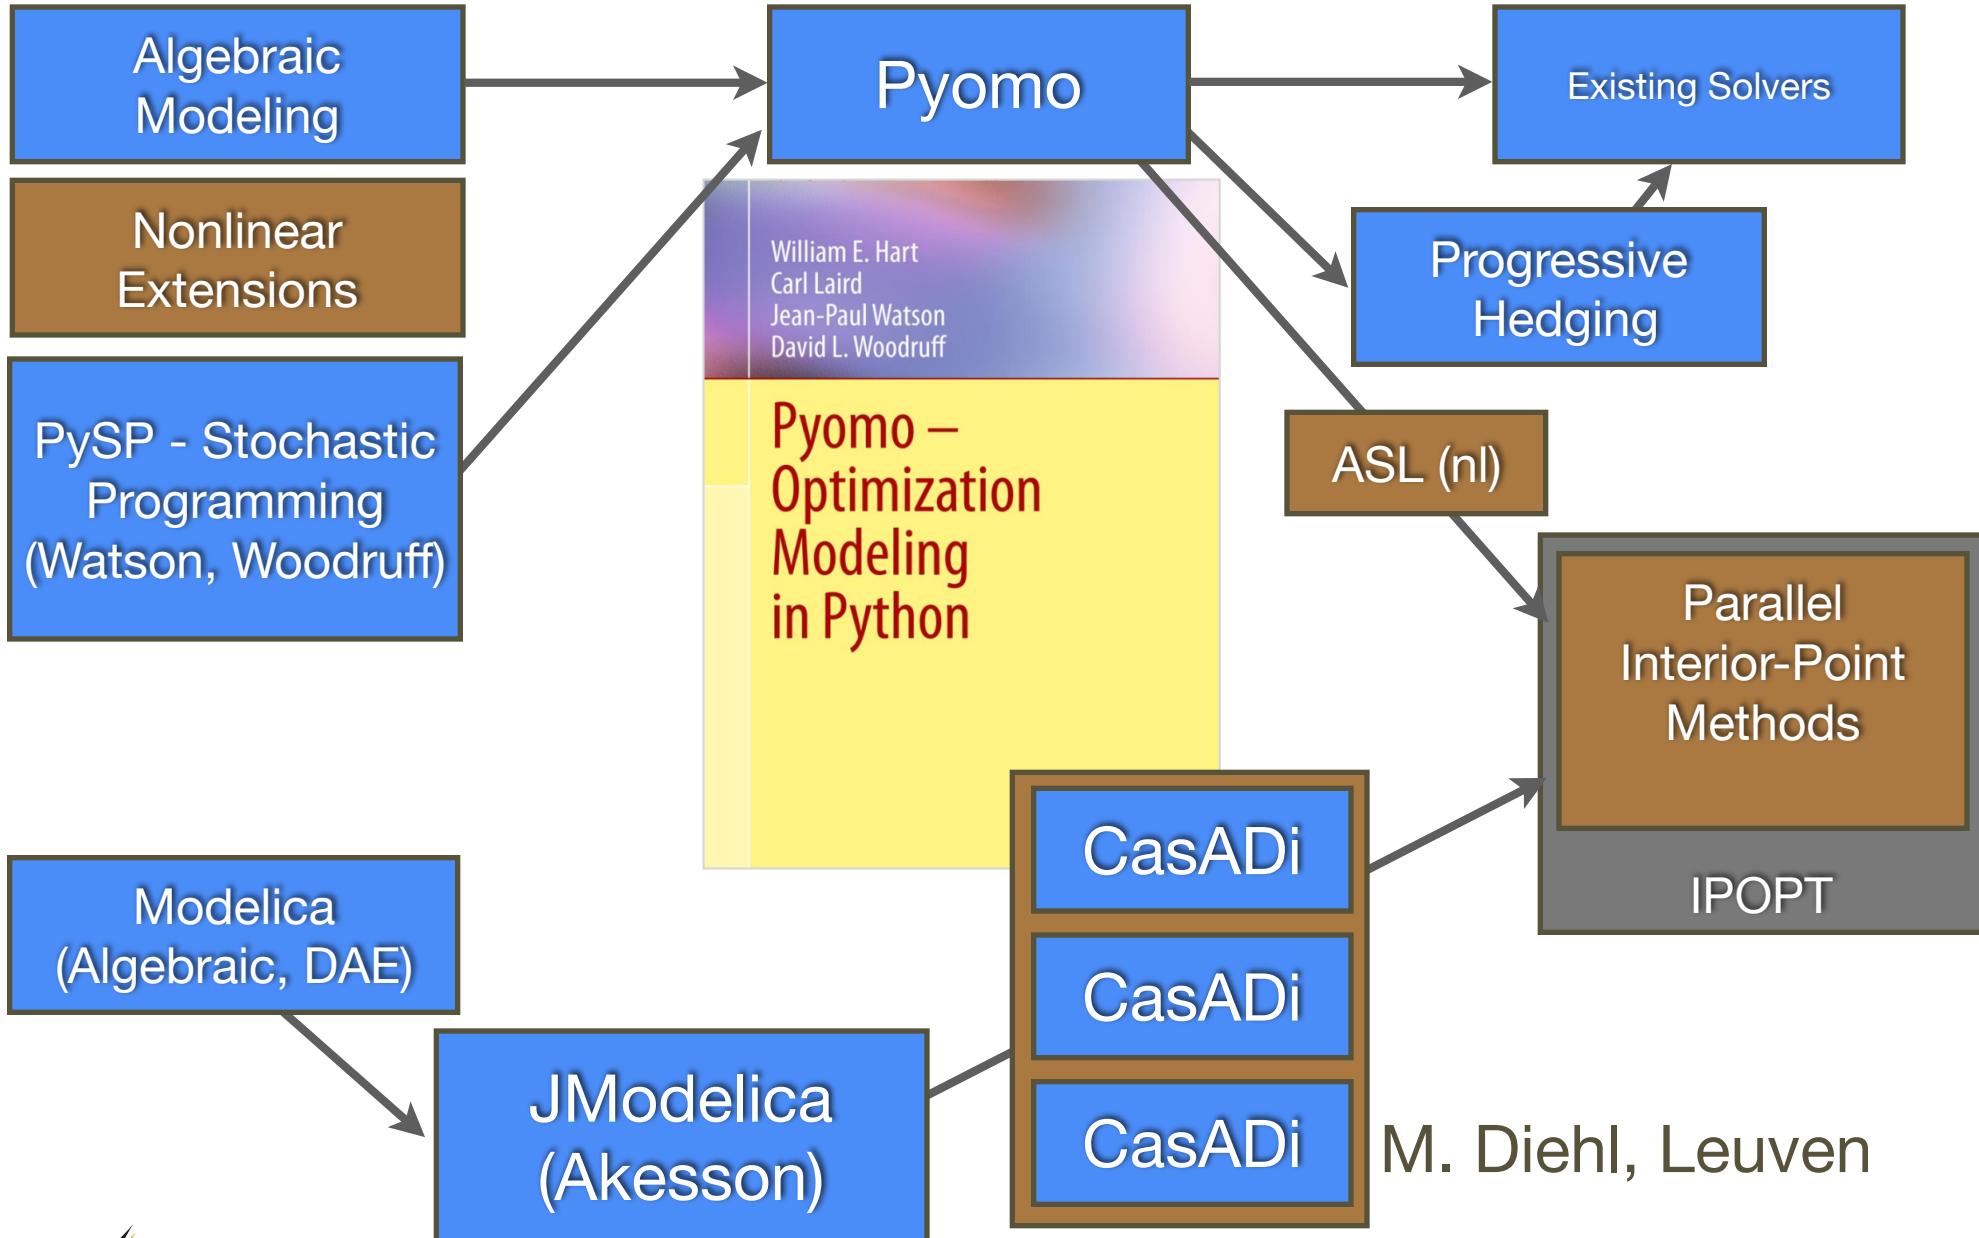
Strong Scaling: Dynamic Optimization

Combined Cycle Power Plant with 10 States and 127 Algebraic Constraints



Word, D.P., Kang, J., Akesson, J., and Laird, C.D., "Efficient Parallel Solution of Large-Scale Nonlinear Dynamic Optimization Problems", to appear in Computational Optimization and Applications, 2014.

Modeling Interfaces



Summary and Conclusions

- Demand for large-scale mathematical programming
 - Large-scale optimization advances allow new strategies
 - For continued improvement, parallel algorithms are necessary
- Applications, Architectures, Algorithms, Adoption
 - Need to understand the applications, assumptions
 - Uncertainty, discretization, spatial/network, data
 - Need to understand architectures
 - E.g. Big-Iron clusters, shared-memory multi-core, GPU
 - Not all parallel architectures created equal: strengths and limitations
 - Need for new algorithms
 - Tailored decomposition based on interior-point methods
 - Problem level decomposition strategies
 - Strategies are problem and architecture dependent
 - Need to make tools available for other researchers
 - Open-source parallel algorithms
 - Pyomo: Optimization Modeling in Python

Acknowledgments

- Current Students/Researchers

- **Jia Kang**
- Arpan Seth
- **Yankai Cao**
- Alberto Benavides-Serrano
- Jianfeng Liu
- Michael Bynum
- Todd Zhen

- Former Students/Researchers

- **Yu Zhu**
- Ahmed Rabie
- George Abbott III
- Chen Wang
- Sean Legg
- Daniel Word
- Angelica Wong
- Xiaorui Yu
- **Gabriel Hackebeil**
- Shawn McGee

- Collaborators

- D. Cummings - JHSPH
- S. Iamsirithaworn - MPHT
- W. Hart, S. McKenna, J.P. Watson, K. Klise, John Siiriola - Sandia
- T. Haxton, R. Murray - EPA
- Johan Akesson, Lund University
- Sam Mannan, TAMU

ALL THINGS OPTIMAL

Home Biography Group Publications Presentations Software Courses

Now at Purdue University
School of Chemical Engineering

Our research group has moved to the School of Chemical Engineering at Purdue University. Please see new contact information below.

All Things Optimal showcases the teaching, research, and service efforts from Carl Laird's research group in the School of Chemical Engineering at Purdue University.

Support

- National Science Foundation Cyber-Enabled Discovery and Innovation (CDI)-Type II
- National Science Foundation (CAREER Grant CBET# 0955205).
- Sandia National Laboratories, EPA, PUB Singapore
- MKOPSC, P2SAC